

# Credit Card fraud detection

The dataset has credit card transactions, and its features are the result of PCA analysis. It has 'Amount', 'Time', and 'Class' features where 'Amount' shows the monetary value of every transaction, 'Time' shows the seconds elapsed between the first and the respective transaction, and 'Class' shows whether a transaction is legit or not.

In 'Class', value 1 represents a fraud transaction, and value 0 represents a valid transaction.

Credit card fraud detection is a critical application of machine learning in finance. It involves a series of steps starting from data collection and preprocessing to model building, evaluation, deployment, and continuous monitoring to safeguard against fraudulent activities.

Overview of the Dataset:

Dataset Source: The dataset consists of credit card transactions made in September 2013 by European cardholders. It contains numerical input variables that are the result of PCA transformations due to privacy concerns. The features include time, amount, and anonymized numerical features (V1-V28) that are a result of PCA transformation.

Objective: The primary objective is to detect fraudulent transactions among a vast number of legitimate ones.

Imbalanced Data: Typically, the dataset is highly imbalanced, where fraudulent transactions are a tiny fraction of the total transactions, making it challenging to train models effectively.

```
In [ ]:
In [ ]: Data Preprocessing: Explore and understand the data. Handle missing values, outliers, and scale/normalize numerical features.
Feature Engineering: Create or extract new features that might aid in fraud detection.
Handling Imbalanced Data: Techniques like oversampling (SMOTE), undersampling, or using algorithms robust to class imbalance.
Model Selection: Commonly used models include Logistic Regression, Decision Trees, Random Forests, Gradient Boosting, and Support Vector Machines.
Model Evaluation: Metrics like precision, recall, F1-score, and area under the ROC curve (AUC-ROC) are used to evaluate model performance.
Hyperparameter Tuning: Optimize model parameters to enhance performance.
Deployment and Monitoring: Deploy the model in a production environment and continually monitor its performance.
```

## Import Models

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')
%matplotlib inline

In [2]: df = pd.read_csv('creditcard.csv')
df.head()
```

Out[2]:	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	V23
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.018307	0.277838	-0.110474
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771679	0.909412
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458

5 rows × 31 columns

```
In [3]: # statistical info
df.describe()
```

Out[3]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	94813.859575	3.918649e-15	5.682686e-16	-8.761736e-15	2.811118e-15	-1.552103e-15	2.040130e-15	-1.698953e-15	-1.893285e-16
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00	1.380247e+00	1.332271e+00	1.237094e+00	1.194353e+00
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00	-1.137433e+02	-2.616051e+01	-4.355724e+01	-7.321672e+01
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01	-6.915971e-01	-7.682956e-01	-5.540759e-01	-2.086297e-01
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02	-5.433583e-02	-2.741871e-01	4.010308e-02	2.235804e-02
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01	6.119264e-01	3.985649e-01	5.704361e-01	3.273459e-01
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01	3.480167e+01	7.330163e+01	1.205895e+02	2.000721e+01

8 rows × 31 columns

In [4]:

```
# datatype info
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Time    284807 non-null   float64
1   V1      284807 non-null   float64
2   V2      284807 non-null   float64
3   V3      284807 non-null   float64
4   V4      284807 non-null   float64
5   V5      284807 non-null   float64
6   V6      284807 non-null   float64
7   V7      284807 non-null   float64
8   V8      284807 non-null   float64
9   V9      284807 non-null   float64
10  V10     284807 non-null   float64
11  V11     284807 non-null   float64
12  V12     284807 non-null   float64
13  V13     284807 non-null   float64
14  V14     284807 non-null   float64
15  V15     284807 non-null   float64
16  V16     284807 non-null   float64
17  V17     284807 non-null   float64
18  V18     284807 non-null   float64
19  V19     284807 non-null   float64
20  V20     284807 non-null   float64
21  V21     284807 non-null   float64
22  V22     284807 non-null   float64
23  V23     284807 non-null   float64
24  V24     284807 non-null   float64
25  V25     284807 non-null   float64
26  V26     284807 non-null   float64
27  V27     284807 non-null   float64
28  V28     284807 non-null   float64
29  Amount  284807 non-null   float64
30  Class   284807 non-null   int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
```

## Preprocessing the Data

In [5]:


```
# check for null values
df.isnull().sum()
```

```
Out[5]: Time      0
V1          0
V2          0
V3          0
V4          0
V5          0
V6          0
V7          0
V8          0
V9          0
V10         0
V11         0
V12         0
V13         0
V14         0
V15         0
V16         0
V17         0
V18         0
V19         0
V20         0
V21         0
V22         0
V23         0
V24         0
V25         0
V26         0
V27         0
V28         0
Amount      0
Class       0
dtype: int64
```

## Exploratory Data Analysis

```
In [6]: sns.countplot(df['Class'])
```


```
Out[6]: <AxesSubplot:xlabel='Class', ylabel='count'>
```

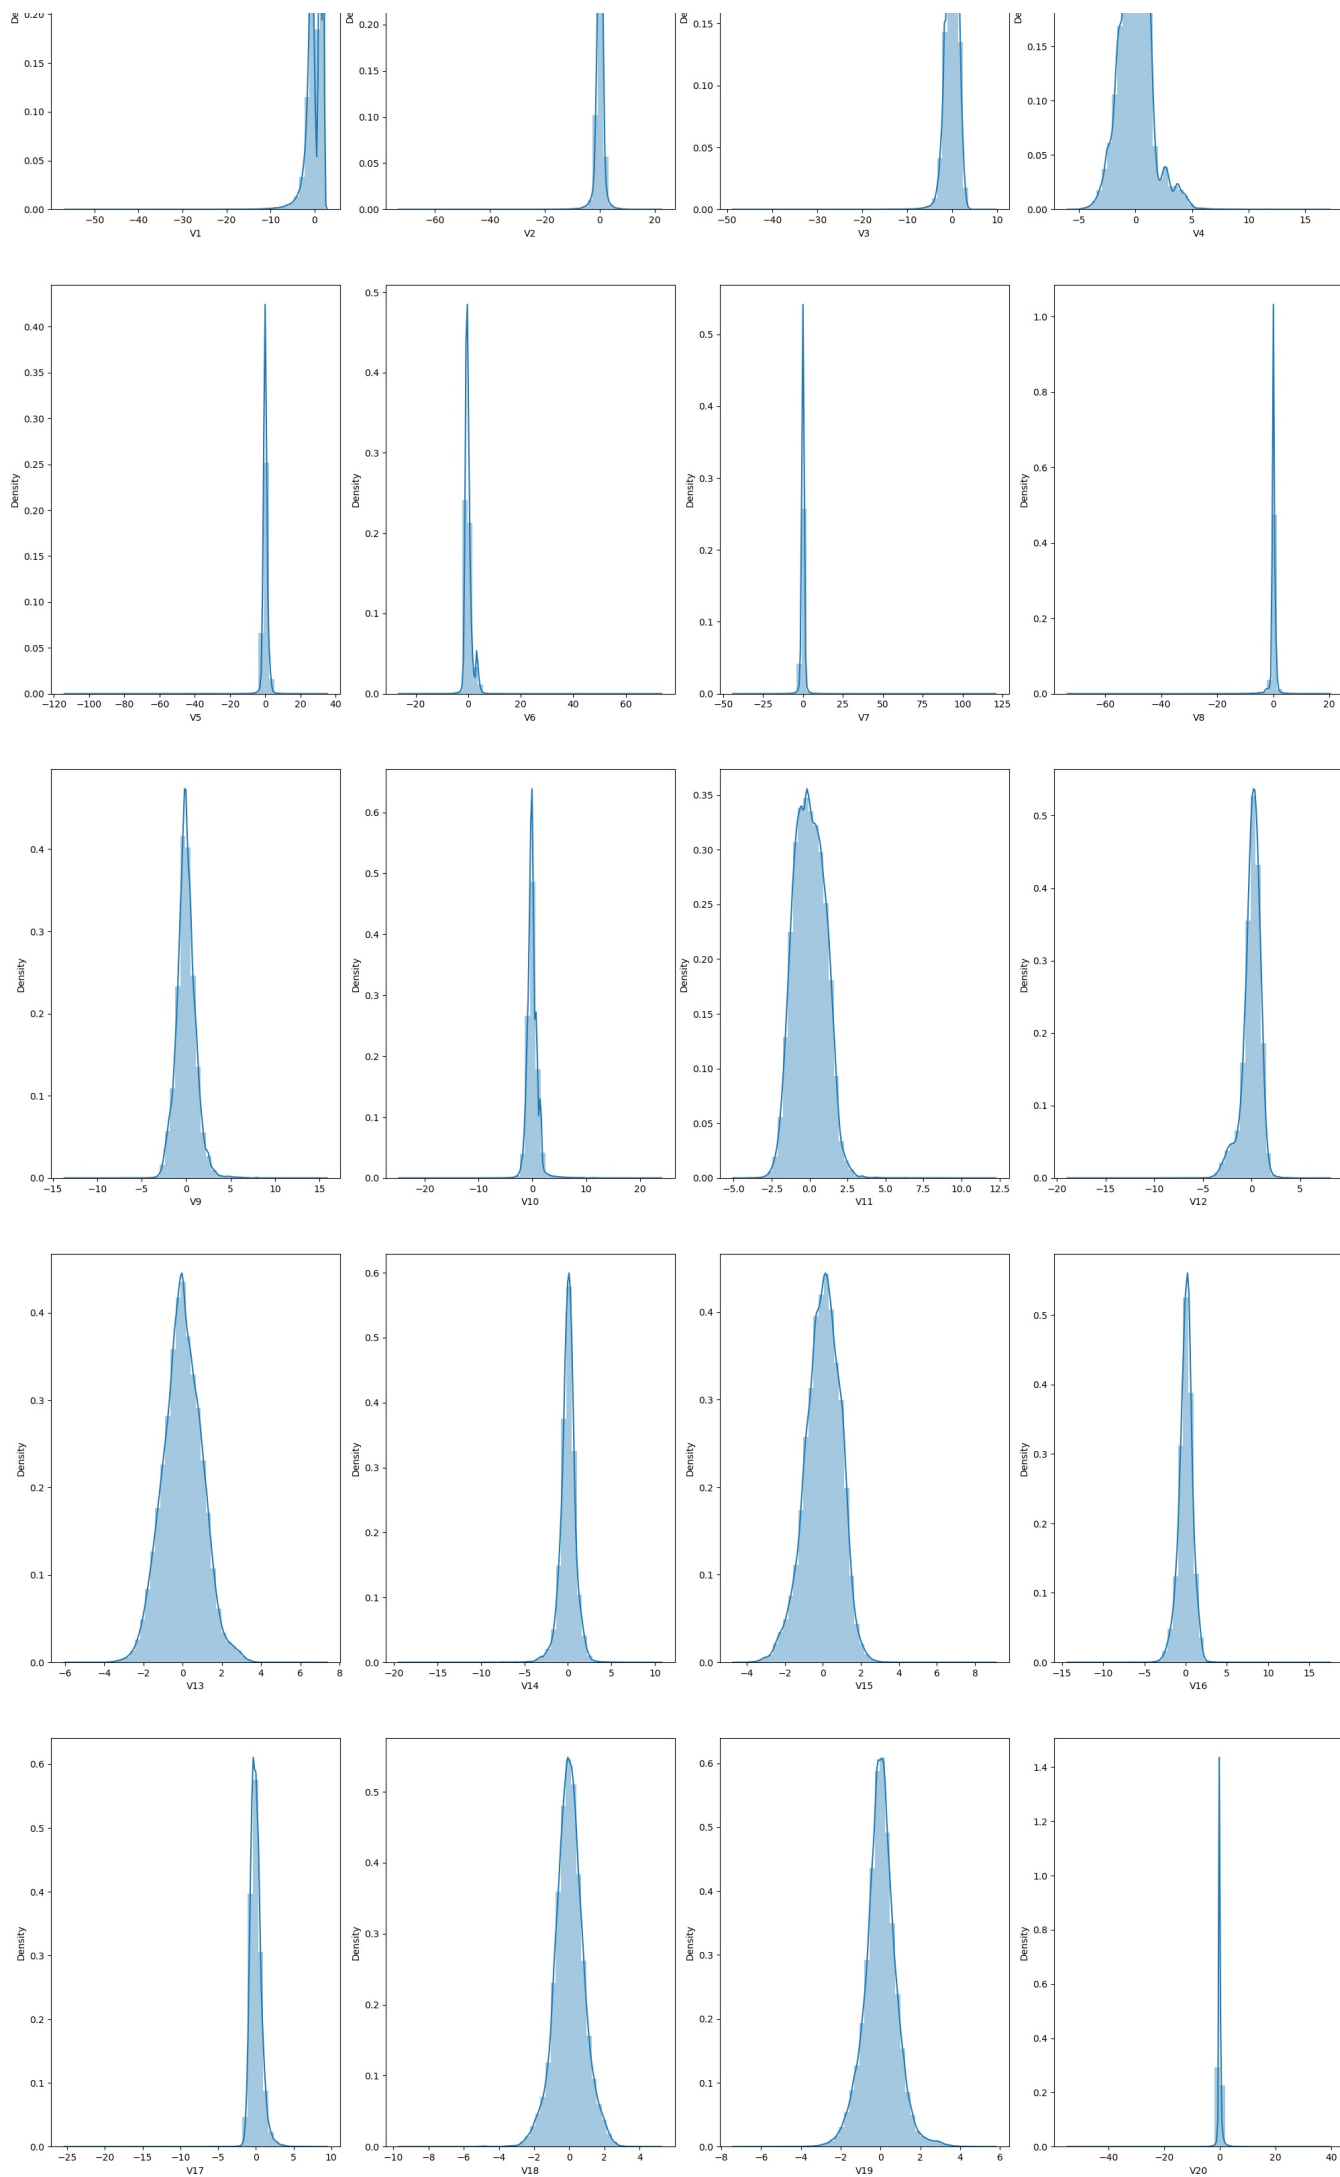



```
In [7]: df_temp = df.drop(columns=['Time', 'Amount', 'Class'], axis=1)
```

```
# create dist plots
fig, ax = plt.subplots(ncols=4, nrows=7, figsize=(20, 50))
index = 0
ax = ax.flatten()

for col in df_temp.columns:
    sns.distplot(df_temp[col], ax=ax[index])
    index += 1
plt.tight_layout(pad=0.5, w_pad=0.5, h_pad=5)
```









```
In [8]: sns.distplot(df['Time'])
```

```
Out[8]: <AxesSubplot:xlabel='Time', ylabel='Density'>
```



```
In [9]: sns.distplot(df['Amount'])
```


```
Out[9]: <AxesSubplot:xlabel='Amount', ylabel='Density'>
```



## Coorelation Matrix

```
In [10]: corr = df.corr()  
plt.figure(figsize=(30,40))  
sns.heatmap(corr, annot=True, cmap='coolwarm')
```

```
Out[10]: <AxesSubplot:>
```



## Input split

```
In [11]: X = df.drop(columns=['Class'], axis=1)
y = df['Class']
```

## Standard scaling

```
In [12]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
x_scaler = sc.fit_transform(X)
```

```
In [13]: x_scaler[-1]
```

```
Out[13]: array([ 1.64205773, -0.27233093, -0.11489898,  0.46386564, -0.35757
, -0.00908946, -0.48760183,  1.27476937, -0.3471764,  0.44253246,
, -0.84072963, -1.01934641, -0.0315383, -0.18898634, -0.08795849,
, 0.04515766, -0.34535763, -0.77752147,  0.1997554, -0.31462479,
, 0.49673933,  0.35541083,  0.8861488,  0.6033653,  0.01452561,
, -0.90863123, -1.69685342, -0.00598394,  0.04134999,  0.51435531])
```

## Model Training

```
In [14]: # train test split
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, f1_score
x_train, x_test, y_train, y_test = train_test_split(x_scaler, y, test_size=0.25, random_state=42, stratify=y)
```

```
In [15]: from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
# training
model.fit(x_train, y_train)
# testing
y_pred = model.predict(x_test)
print(classification_report(y_test, y_pred))
print("F1 Score:", f1_score(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	71079
1	0.85	0.63	0.72	123
accuracy			1.00	71202
macro avg	0.92	0.81	0.86	71202
weighted avg	1.00	1.00	1.00	71202

F1 Score: 0.719626168224299

```
In [20]: from sklearn.ensemble import RandomForestClassifier
model = RandomForestClassifier()
# training
model.fit(x_train, y_train)
# testing
y_pred = model.predict(x_test)
print(classification_report(y_test, y_pred))
print("F1 Score:", f1_score(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	71079
1	0.96	0.78	0.86	123
accuracy			1.00	71202
macro avg	0.98	0.89	0.93	71202
weighted avg	1.00	1.00	1.00	71202

F1 Score: 0.8609865470852018

```
In [17]: from xgboost import XGBClassifier
model = XGBClassifier(n_jobs=-1)
# training
model.fit(x_train, y_train)
# testing
y_pred = model.predict(x_test)
print(classification_report(y_test, y_pred))
print("F1 Score:", f1_score(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	71079
1	0.94	0.79	0.86	123
accuracy			1.00	71202
macro avg	0.97	0.89	0.93	71202
weighted avg	1.00	1.00	1.00	71202


F1 Score: 0.8584070796460177

## Class Imbalancement

```
In [18]: sns.countplot(y_train)
```

```
Out[18]: <AxesSubplot:xlabel='Class', ylabel='count'>
```






```
In [21]: # balance the class with equal distribution
from imblearn.over_sampling import SMOTE
over_sample = SMOTE()
x_smote, y_smote = over_sample.fit_resample(x_train, y_train)
```

```
In [22]: sns.countplot(y_smote)
```

```
Out[22]: <AxesSubplot:xlabel='Class', ylabel='count'>
```



```
In [23]: from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
# training
model.fit(x_smote, y_smote)
# testing
y_pred = model.predict(x_test)
print(classification_report(y_test, y_pred))
print("F1 Score:", f1_score(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	0.98	0.99	71079
1	0.06	0.89	0.11	123
accuracy			0.98	71202
macro avg	0.53	0.93	0.55	71202
weighted avg	1.00	0.98	0.99	71202

F1 Score: 0.11219763252702007

```
In [24]: from sklearn.ensemble import RandomForestClassifier
model = RandomForestClassifier(n_jobs=-1)
# training
```

```
model.fit(x_smote, y_smote)
# testing
y_pred = model.predict(x_test)
print(classification_report(y_test, y_pred))
print("F1 Score:", f1_score(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	71079
1	0.88	0.77	0.82	123
accuracy			1.00	71202
macro avg	0.94	0.89	0.91	71202
weighted avg	1.00	1.00	1.00	71202

F1 Score: 0.8225108225108225

In [ ]: